

The agent readiness test

Ten questions to answer before you let an agent run. Without clear answers, stay with a project or a manual workflow for now. That is good system design.

Agent Owner Date

01 What single recurring problem does it solve?

02 What triggers it?

03 What are the exact steps?

04 Which sources can it trust?

05 Which tools does it need?

06 Which actions can it take without approval?

07 Which actions always require approval?

08 What must it never do?

09 Who reviews the result?

10 How will we know it is working?

Where not to start

THE BEST FIRST AGENT IS USUALLY BORING. A BORING WORKFLOW IS EASIER TO UNDERSTAND, TEST AND IMPROVE.

- x** Unsupervised external political communication
- x** Contacting stakeholders in your name
- x** Changing organisational positions
- x** Publishing social content
- x** Making sensitive strategic decisions
- x** Interpreting ambiguous political developments without review
- x** Maintaining sensitive personal profiles
- x** Taking irreversible actions

The build checklist

Tick each line before the agent runs on a schedule. Anything unticked is a decision you have not made yet.

Instructions

- ☐ Mission is specific.
- ☐ Role is defined.
- ☐ Audience is defined.
- ☐ Workflow is written step by step.
- ☐ Output format is defined.
- ☐ Uncertainty rules are defined.

Boundaries

- ☐ Prohibited actions are explicit.
- ☐ Human approval points are defined.
- ☐ Sensitive situations trigger escalation.
- ☐ Failure behaviour is defined.
- ☐ Recipient lists are controlled.

Knowledge

- ☐ Sources are trusted.
- ☐ Old material has been reviewed.
- ☐ Rejected drafts are excluded.
- ☐ Important primary sources are included.
- ☐ AI-generated material is checked before it counts as fact.

Testing

- ☐ Workflow has been run manually.
- ☐ Outputs have been reviewed.
- ☐ Instructions have been revised.
- ☐ Sources have been checked.
- ☐ Scheduling comes only after the workflow is stable.

Capabilities

- ☐ Every connected tool has a clear purpose.
- ☐ Permissions are no broader than necessary.
- ☐ External actions are limited.
- ☐ Internal and external actions are distinguished.

Twelve mistakes to check yourself against

- Building an agent before documenting the existing workflow.
- Giving it a vague role.
- Treating instructions as a one-line prompt.
- Feeding it unreviewed AI-generated knowledge.
- Uploading every document available instead of selecting trustworthy material.
- Giving it more capabilities than the job requires.
- Forgetting to define what it cannot do.
- Automating external communication too early.
- Failing to name a human owner.
- Leaving out when the agent should stop or escalate.
- Scheduling a broken workflow and letting it fail repeatedly.
- Assuming autonomy removes the need for verification.

What still has to be decided

You design the system. The agent runs it.